
VNX TA Design, SplunkWeb and Python

Ken Chen
2014/08/07

Agenda

- Python Overview
- SplunkWeb
 - The architecture behind
 - How everything hook up together
- VNX TA
 - Current design
 - Testing
 - Postmortem
- Lessons learned

Life is short,
I use Python

Python is a perfect language!

- either you are a fool or
- a sales man

— Ken

C++ is a perfect language!

- either you are a fool or
- a sales man

— Bjarne

Python Overview



Python Overview

- Python in one slide
 - A (backend) script language
 - Auto memory mgmt by reference counting
 - A Dynamic typing language
 - A Strongly typed language
 - Object oriented/based
 - Platform independent

Python Overview

- Python in one slide (I am lying)
 - One language spec, different implementation
 - CPython (C)
 - Jython (Java)
 - Cython (Written in Python, translate to C)
 - PyPy (RPython/Written in Python/GIT)
 - IronPython (.NET)
 - ...

Python Overview

- The good
 - Simple
 - One coding style (important)
 - One and the only best way to solve a problem
 - **VS** 16 ways to stack a cat (Bjarne)
 - Has good language features
 - Rich standard and 3rd party libraries
 - Community is super good
 - Dynamic typing
 - Productive

Python Overview

- The bad
 - Generally slow (We will see)
 - Single thread due to GIL (We will see)
 - Py3k is not backward compatible with Python 2

Python Overview

- How Slow
 - Let's do a file system scanning
 - Let's do a Fibonacci
- Why slow
 - 90 % of the time, we program in a slow way
 - function call is relative expensive in Python
 - The missing high performance APIs
 - Not pythonic

Python Overview

- The missing APIs
 - `def unmarshal(filename, ObjectClass):`
 `objs = []`
 `with open(filename) as f:`
 `for lin in f:`
 `lin = lin.strip()`
 `if lin:`
 `objs.append(ObjectClass(lin))`
 `return objs`

Python Overview

- Pythonic (Python idiom)
 - Code style
 - “if file == None” is not Pythonic
 - Performance
 - “for i in range(100000)” is not Pythonic
 - There are a few guidelines, but far less than the the guidelines in C++
 - Am I talking about micro-optimisation ?

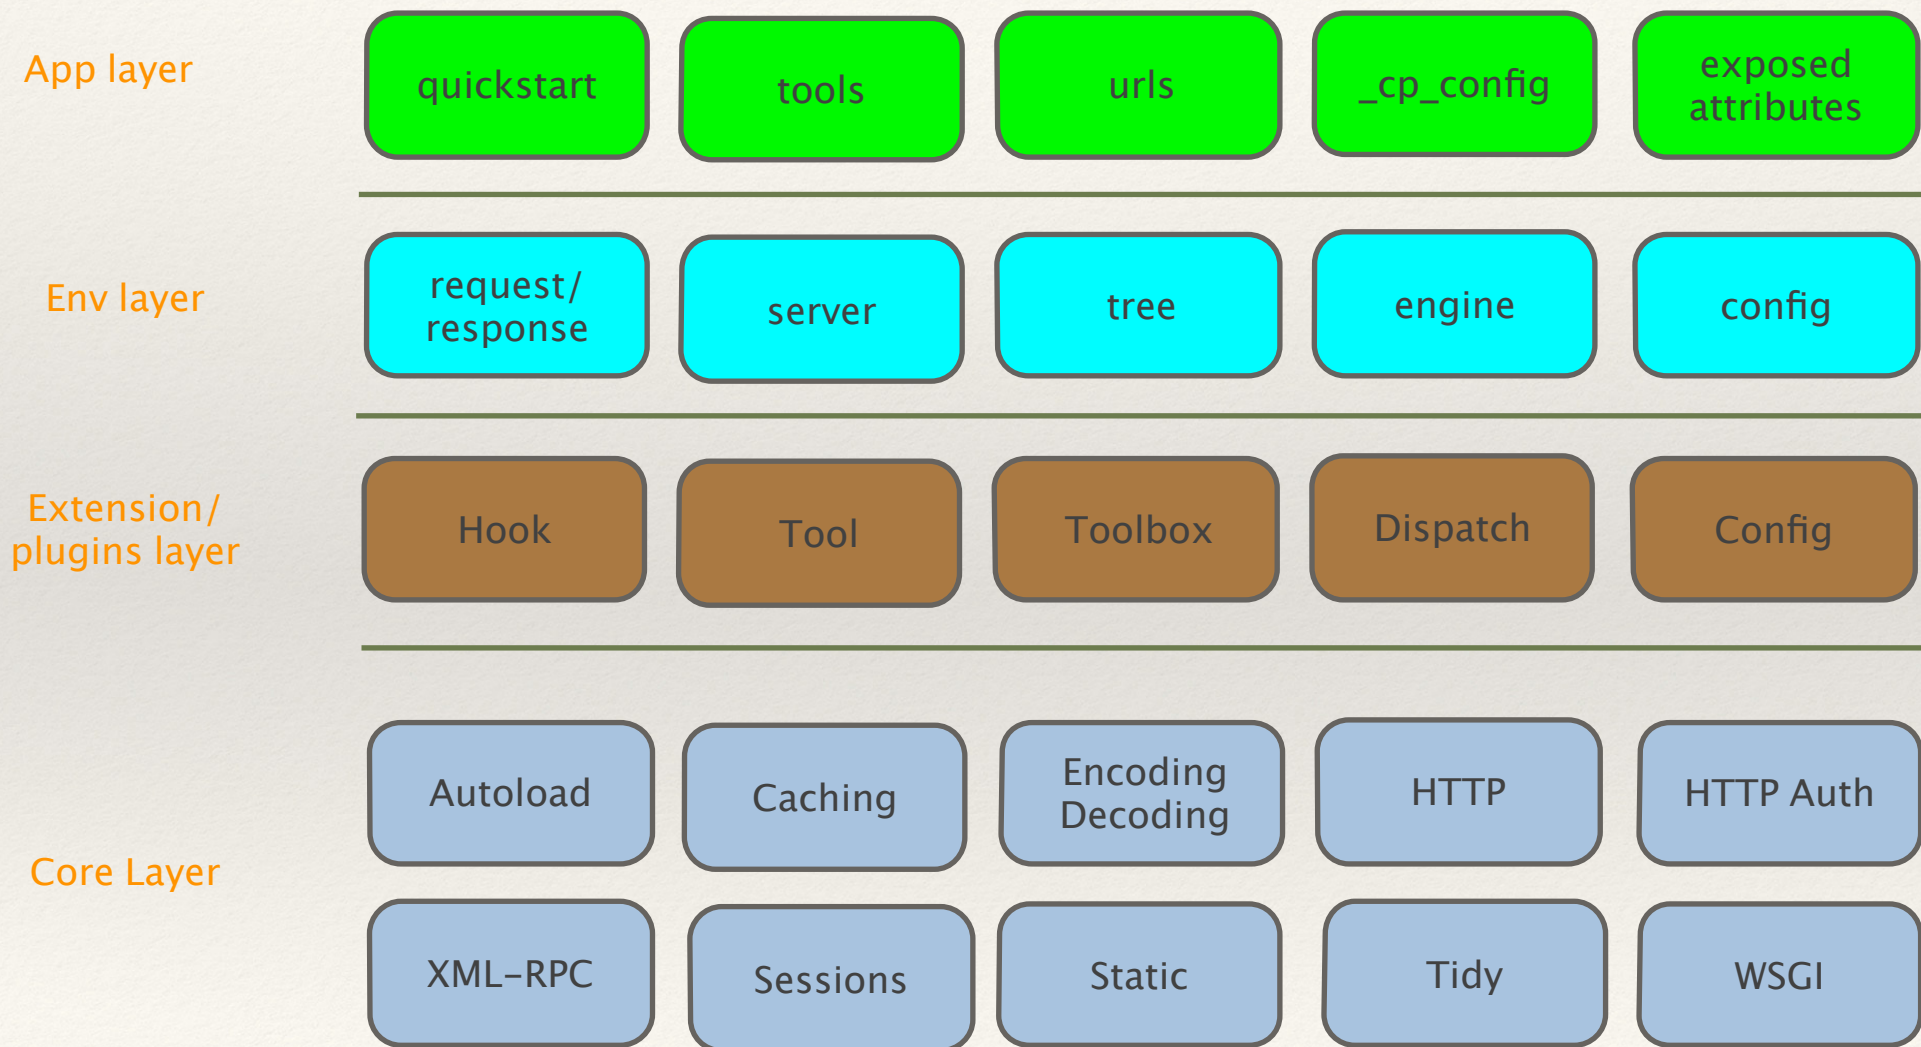
Python Overview

- How to stay Pythonic
 - Read code (Python standard lib)
 - VIM + Plugins (Allow me to show you)
 - pep8.py
 - pylint
 - PyCharm or other IDE

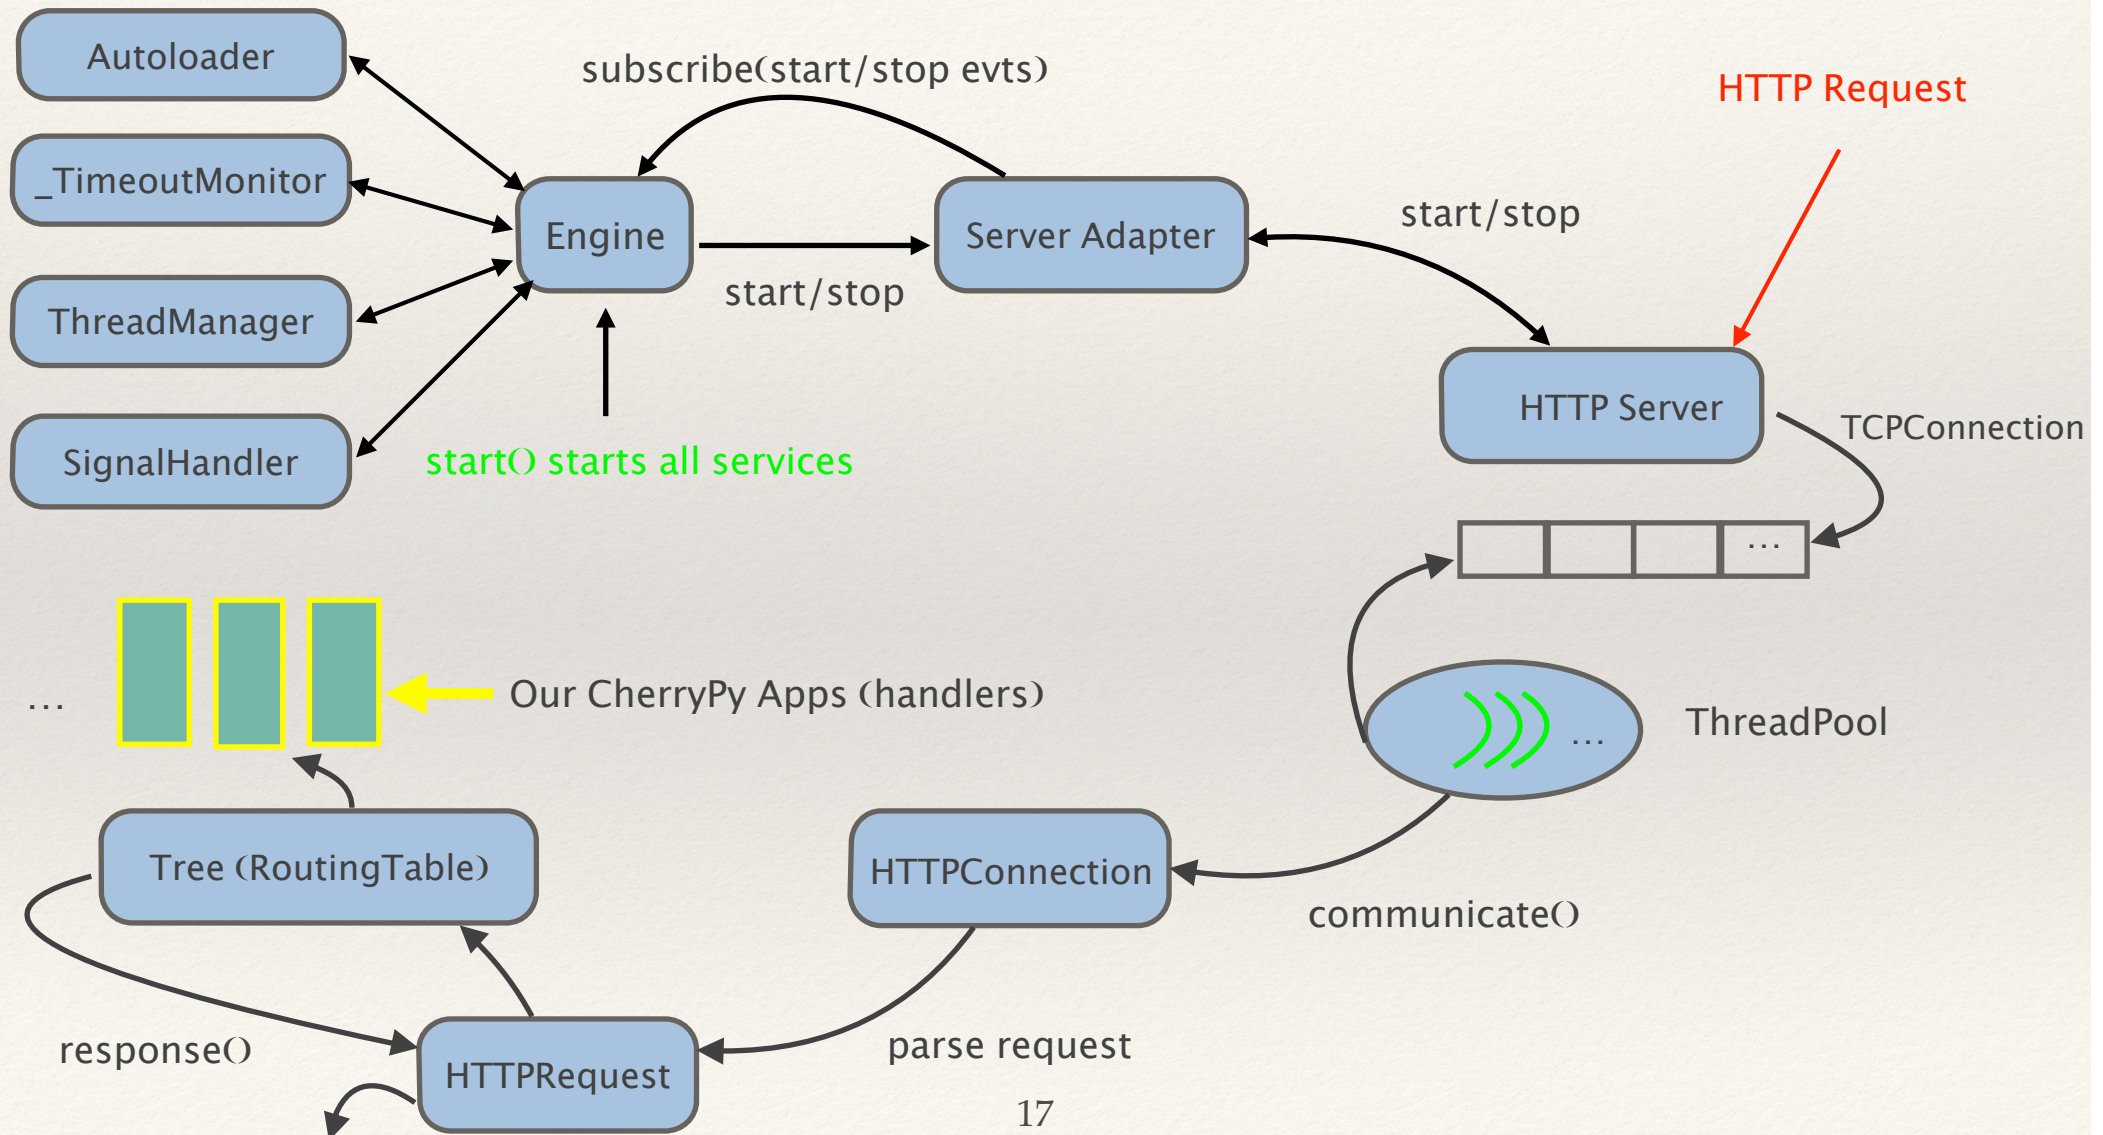
Splunkweb server



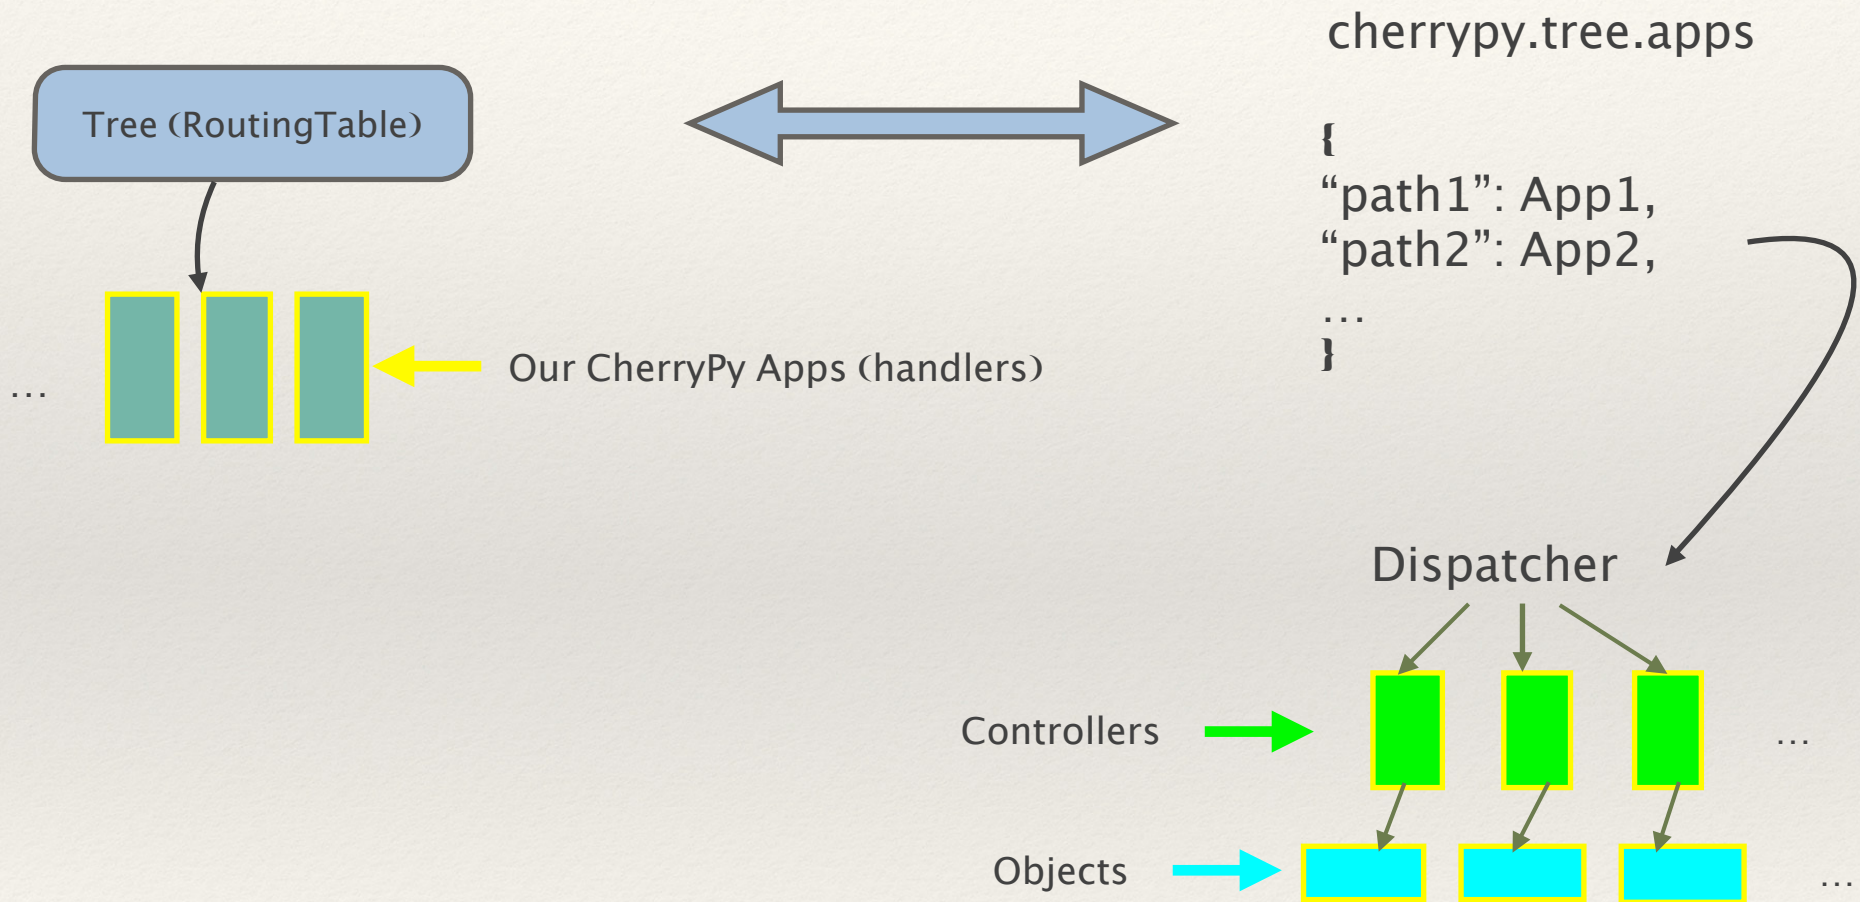
CherryPy



CherryPy architecture



CherryPy the routing table

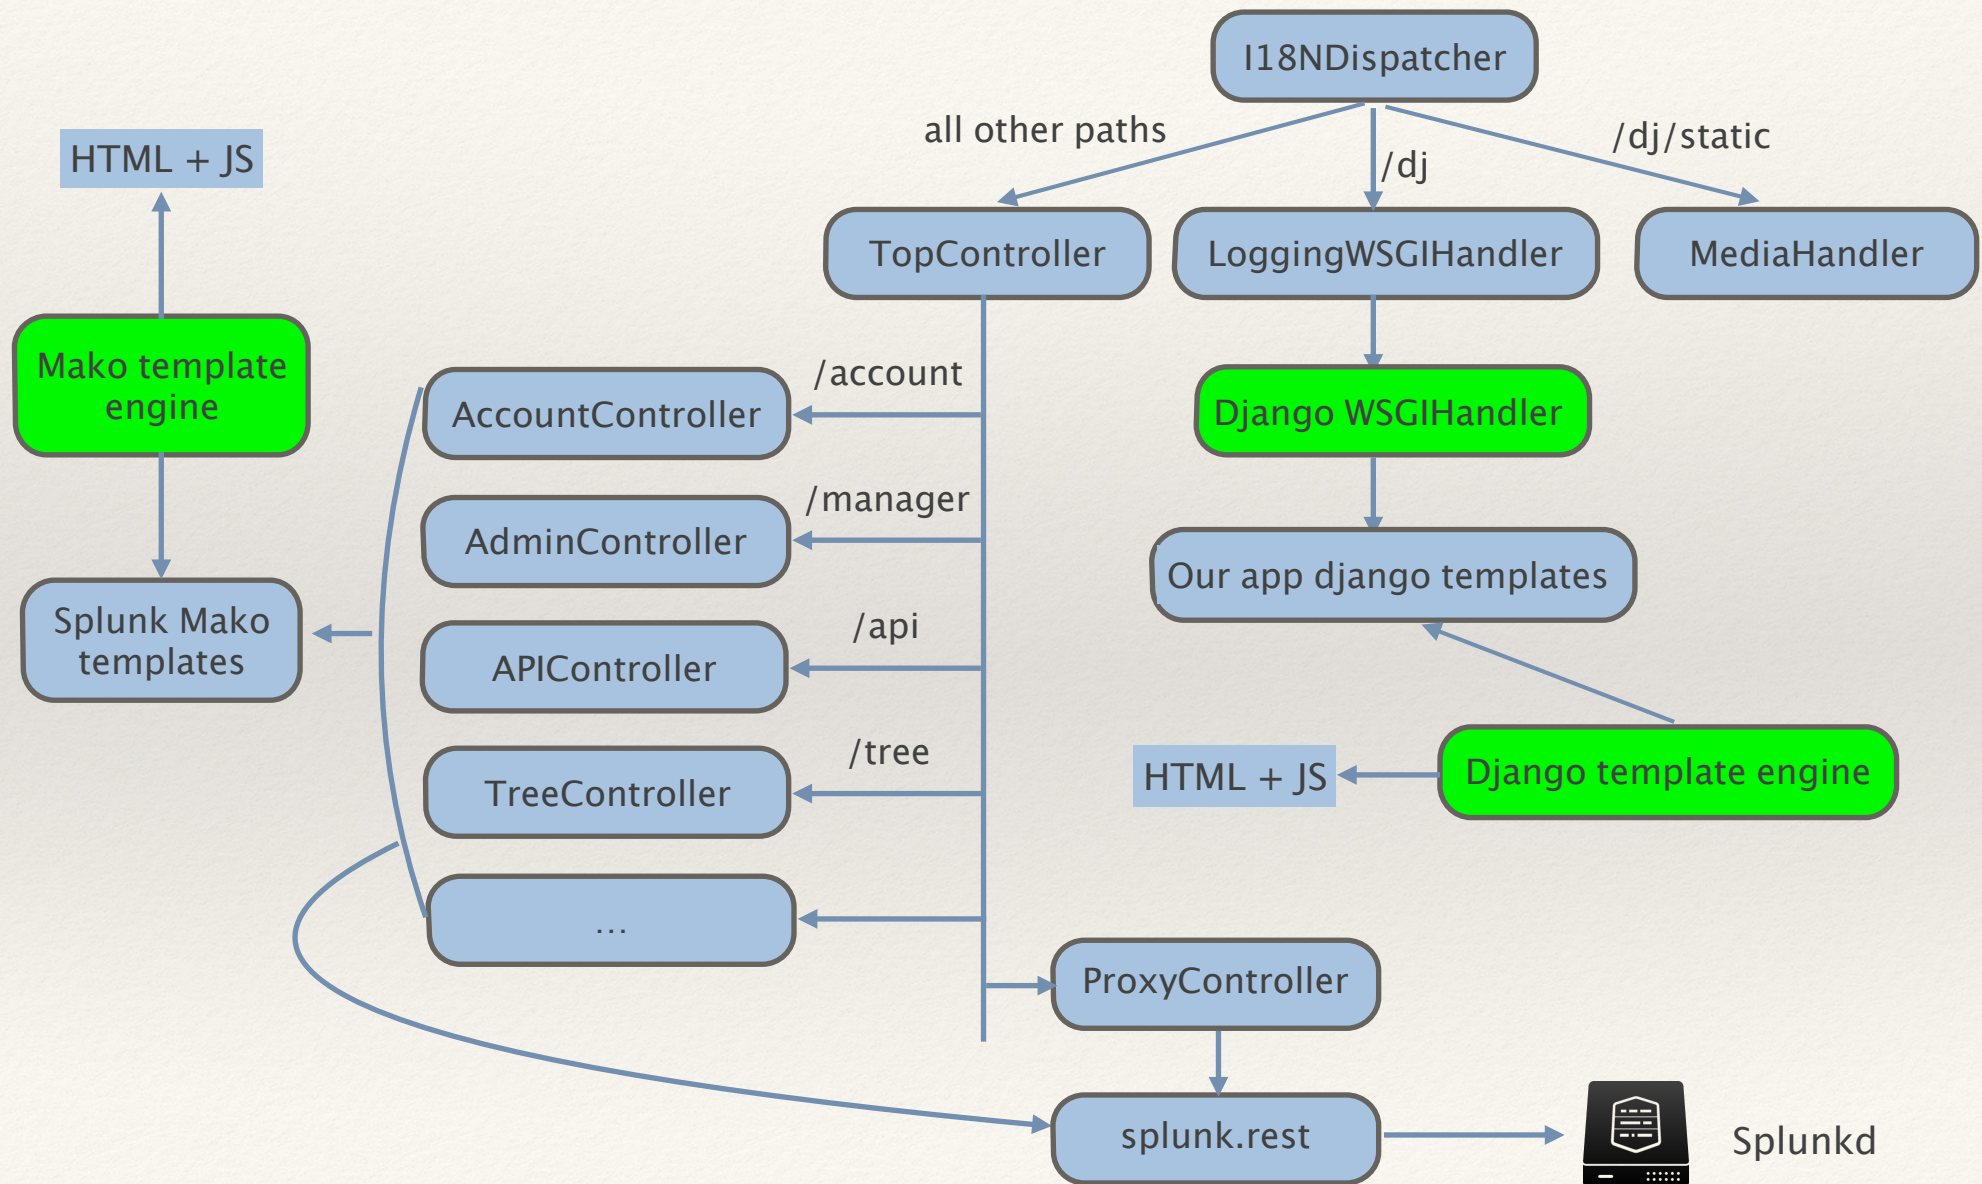


Splunk Apps (Django/CherryPy)

Glue up everything

- Discover CherryPy handlers
 - ControllerMetaLoader discovers all apps by matching “**etc/apps/<app>/appserver/controllers**” folder
 - CustomController attaches handler’s endpoints
 - app’s CherryPy controller inherits from “BaseController”
- Discover Django apps
 - \$SPLUNK_HOME/etc/apps/framework/server/splunkdj/loaders/apps_finder.py
 - Matching “**etc/apps/<app>/django/<app>**” folder

The routing tree



SplunkWeb backend in 2 slides

- The third parties
 - CherryPy
 - Django
 - Mako lib
 - OpenSSL
 - httplib2
 - lxml
 - Others – Babel, beaker, reportlab

SplunkWeb backend in 2 slides

- In house components
 - App Controllers
 - clilib
 - Django templates/tags
 - Bridge Django CORE and Django App
 - Mako templates

VNX TA Design



VNX TA main functionalities

- Data collection
 - VNX File, inventory/perf metrics
 - VNX Block, inventory/status/perf
 - Output parsing
 - Storage objects unmarshaling
- CIM Mapping
- Eventgen sampling

VNX TA main functionalities


- Key challenges
 - VNX itself is hard hard to manage, not a unified system, no unified APIs
 - Multiple machines
 - Blocking call and many calls
 - VNX mgmt path is not designed for high concurrency and perf
 - Python language limitation (GIL)
 - Testing env (\$\$\$)

GIL


- The Global Interpreter Lock
 - Only one thread is able to run
 - Parallel executing the code is not possible
 - Concurrency is still possible
 - Concurrency is not parallelism (Rob Pike, Go)

GIL

- CPU Bound



- IO Bound



GIL

- Python/ceval.c (CPython 2.7.8)

```
if (--_Py_Ticker < 0) {
```

Counting down

```
...
```

```
    _Py_Ticker = _Py_CheckInterval;
```

Reset the counter

```
...
```

```
    if (pendingcalls_to_do) {
```

```
        Py_MakePendingCalls();
```

```
        ...
```

```
    }
```

Handle signals

```
if (interpreter_lock) {
```

```
    /* Give another thread a chance */
```

```
...
```

```
    PyThread_release_lock(interpreter_lock);
```

```
    /* Other threads may run now */
```

```
    PyThread_acquire_lock(interpreter_lock, 1);
```

```
...
```

```
    /* Check for thread interrupts */
```

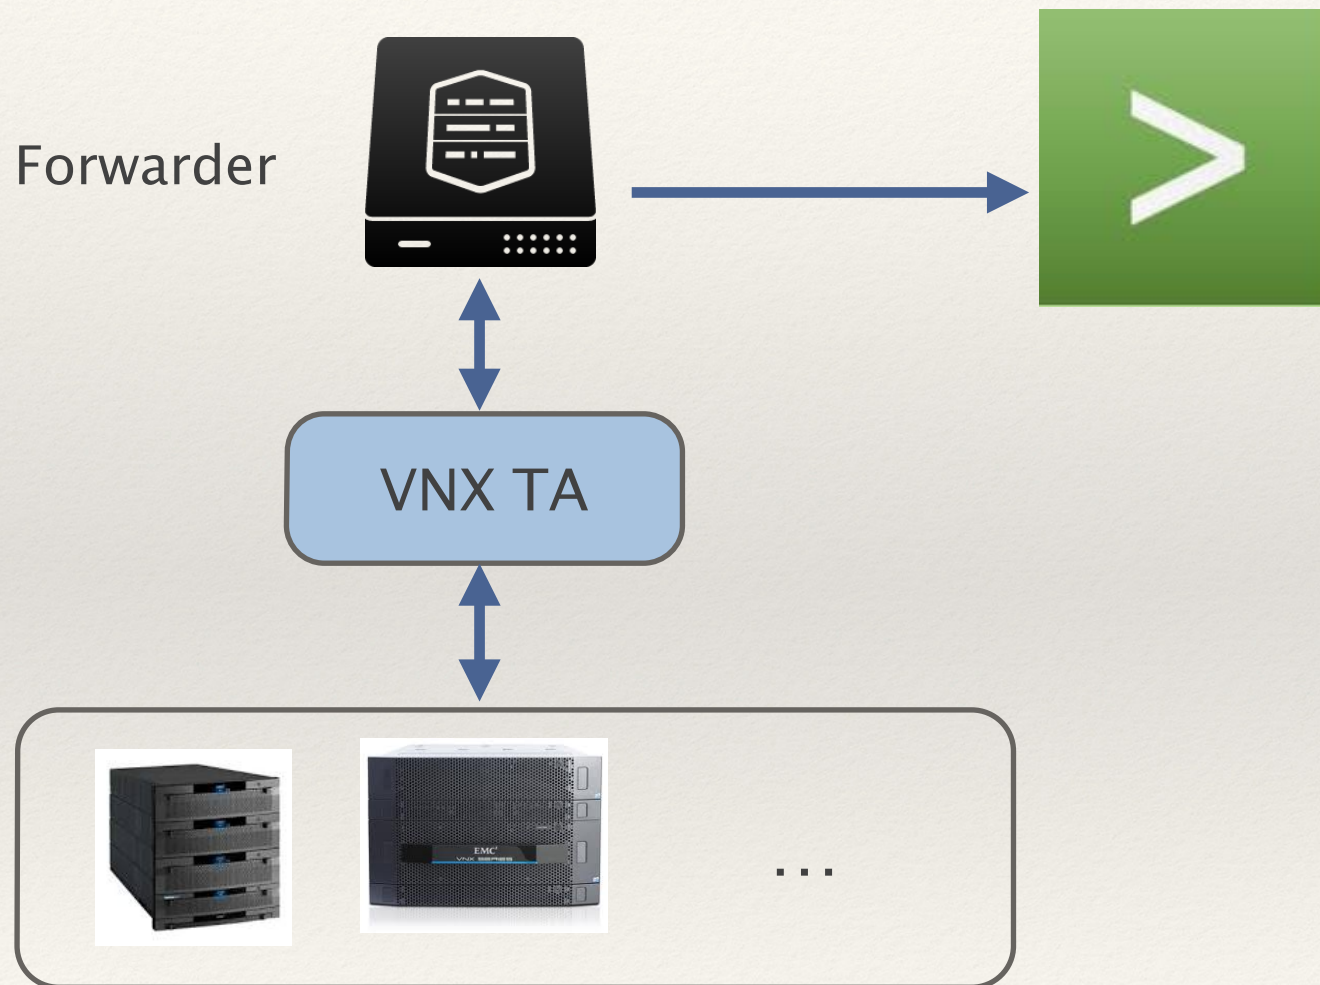
```
...
```

```
    }
```

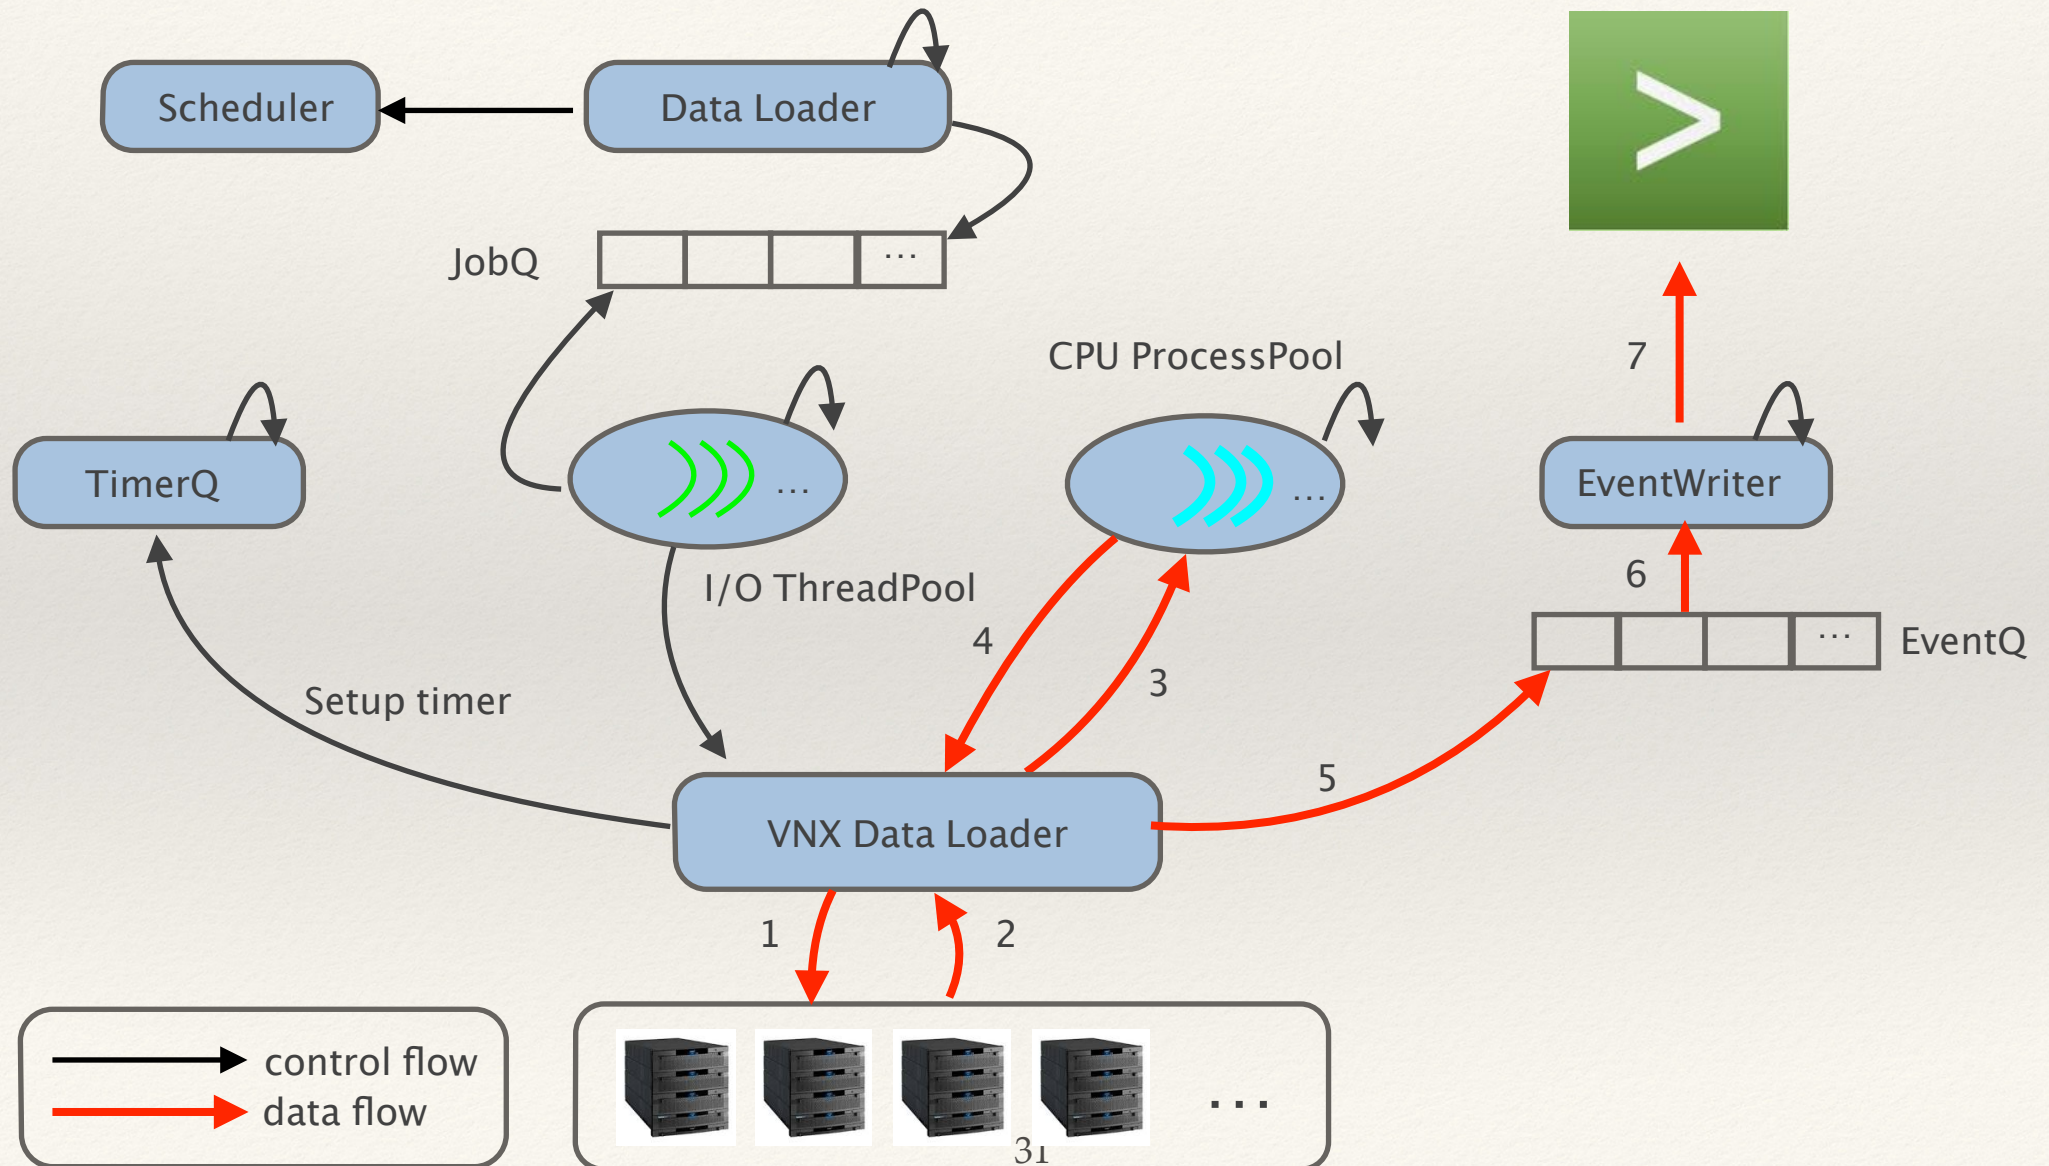
```
}
```

Release and reacquire the lock

VNX TA high level design



VNX TA architecture



Pros and cons

- Pros
 - Long lived, no startup overhead
 - Gracefully handle network/load jitter
 - No checkpoint. Maintain status in memory
 - Multithread for I/O, multiprocess for computing, effective workaround GIL
 - Time service handles hung requests
 - Graceful tear down (pay attention to you signal handler)
- Cons
 - SHELL is hard to scale up as splunkd does(async I/O is the right way). Randomization relieves the problem somehow
 - Relatively Complicate

Testing

- Unit test is an absolute need
 - It is hard to test the parser (different output in different version)
 - Need setup test strategy
 - The TA employs replaying outputs from different VNX version
- Performance test
 - Mocks up CLI, simulate verbose output

Postmortem

- Logging
 - Separate the logs (Generic ones and BLOs)
 - Timestamp and thread id is necessary
 - Support debug level switch
 - Verbose logging
 - All raw outputs are captured and saved

Lessons learned

- Embrace GIL
- Embrace threading (I/O)
- Embrace multiprocessing (computation)
- Embrace pythonic
- Embrace pdb (cover little in this session)
- Embrace logging
- Program at high level and think at low level

Life is short,
You need Python

Thread termination

Destroy me !